

# Assignment 1: Lexical Analysis with Flex

CS 464/5607: Compiler Design

Spring 2026

## 1 Introduction

In this assignment, you will implement the first phase of a compiler: the **Lexical Analyzer** (or Scanner). Your task is to use the **flex** tool to generate a C scanner that reads source code and breaks it down into a stream of tokens.

You will be provided with a driver program and a list of required tokens. Your job is to define the Regular Expressions (Regex) and rules in a `.l` file to correctly identify these tokens.

## 2 Understanding Flex (The Fast Lexical Analyzer Generator)

### 2.1 What is Flex?

Flex (Fast Lexical Analyzer Generator) is a tool used to generate scanners: programs that recognize lexical patterns in text. It reads a specification file (typically with a `.l` extension) containing pairs of Regular Expressions and C code actions, and it outputs a C source file (usually `lex.yy.c`) that implements a scanner for those patterns.

### 2.2 Why are we using it?

Writing a lexical analyzer by hand involves creating a complex Finite State Machine (FSM) to track state transitions for every character. While possible, this is tedious and error-prone. Flex automates this process. By simply describing *what* patterns you want to match using high-level Regular Expressions, Flex handles the difficult task of generating the optimized, low-level C code to implement the corresponding FSM.

### 2.3 How Flex Works

When you run the command `flex lexer.l`, the following happens in the background:

1. **Regex Parsing:** Flex reads your rules and parses the Regular Expressions.
2. **NFA Construction:** It converts each Regex into a Nondeterministic Finite Automaton (NFA).
3. **DFA Conversion:** It combines these NFAs and converts them into a single Deterministic Finite Automaton (DFA), which is much faster to execute.
4. **Table Generation:** It generates transition tables and a driver routine (the `yylex()` function) in C.

The resulting file, `lex.yy.c`, is standard C code. When compiled and linked with your driver program, the `yylex()` function simulates the DFA to tokenize the input stream efficiently.

Note: You will need Flex installed on your system to complete this assignment.

### 3 Provided Files

The following files are provided in the skeleton code. You should not modify the C++ files or headers; your work is focused on the Flex file.

#### 3.1 driver.cpp

This is the entry point of the application. Its purpose is to:

1. Open the input file.
2. Call the `yylex()` function (generated by Flex) in a loop.
3. Print the token type and the matched lexeme to standard output using `std::cout`.

#### 3.2 tokens.h

This header file defines the `enum` for all the Token IDs (e.g., `TOKEN_INT`, `TOKEN_IF`, `TOKEN_ID`). **Purpose:** These IDs are the “return values” your lexer must send back to the driver whenever it recognizes a pattern.

#### 3.3 lexer.l (Your Task)

This is the Flex specification file where you will write your implementation. You must define the patterns (Regex) to recognize keywords, operators, identifiers, and literals.

## 4 Implementation Guide

### 4.1 Structure of a Flex File

A Flex file consists of three distinct sections separated by `%`:

1. **Definitions Section:** Used for C imports and defining reusable Regex macros.
2. **Rules Section:** This is the core of your assignment. Here you define the patterns (using Regex) and the corresponding actions (C code) to execute when a pattern is matched.
3. **User Code Section:** Used for helper C functions (not required for this assignment).

Your goal is to populate the Rules Section to recognize the language constructs.

### 4.2 Implementation Hints

Here are a few tips to help you structure your `lexer.l` file effectively.

#### 4.2.1 Regex Definitions

To keep your rules clean and readable, you should define reusable Regular Expressions in the **Definitions Section** (the top part of the file). This allows you to construct complex patterns from simpler building blocks.

##### Example of Regex Definitions:

```
1 DIGIT      [0-9]
2 LETTER     [a-zA-Z]
```

### 4.2.2 Handling Whitespace and Comments

The lexical analyzer should only return meaningful tokens to the parser.

- **Whitespace:** Spaces, tabs, and newlines should be ignored (consumed without returning a token).
- **Comments:** You must implement support for **single-line comments** (starting with `//`). These should also be ignored. Note that multi-line comments are not supported in this assignment.

### 4.2.3 Token Definitions

Refer to `tokens.h` for the exact names of the token constants you need to return (e.g., `TOKEN_WHILE`, `TOKEN_PLUS`). Your lexer must support all tokens listed in that file.

### 4.2.4 Handling Errors

A robust lexer must handle unexpected input. If the scanner encounters a character or pattern that does not match any of your defined rules, it should return `T_ERROR`. This acts as a catch-all for invalid characters.

To provide meaningful error messages, Flex provides two global variables:

- `yylineno`: An integer holding the current line number (enabled by `%option yylineno`).
- `yytext`: A string containing the text that matched the current rule.

You can use these variables in your error rule to print exactly where and why the error occurred. For example:

```
1 printf("Lexical Error at line %d: Unknown char '%s'\n", yylineno, yytext);
```

**Note on Order:** Since `T_ERROR` is a catch-all, where you place this rule matters significantly. (See Section 3.3).

## 4.3 Key Lexical Properties

To implement the lexer correctly, you must understand how Flex decides which rule to apply when multiple patterns could potentially match the input.

### 4.3.1 1. The Longest Match Rule (Maximal Munch)

Flex will always choose the rule that matches the **most characters** in the input stream.

**Conceptual Example:** Imagine your input is the word "integer".

- You have a rule that matches the keyword "int".
- You have another rule that matches identifiers (words starting with a letter).

Even though the first 3 characters match the keyword rule, Flex sees that the identifier rule can match all 7 characters ("integer"). Because 7 is greater than 3, Flex selects the identifier rule. This ensures that variables with names like `integer` or `format` are not incorrectly split into keywords.

### 4.3.2 2. Rule Priority (Order Matters)

What happens if two rules match the **exact same number** of characters? In this case, Flex resolves the tie by choosing the rule that appears **first** in the `lexer.l` file.

**Conceptual Example:** Imagine your input is the word "int".

- The keyword rule matches 3 characters.
- The identifier rule also matches 3 characters.

If the identifier rule is placed *before* the keyword rule in your file, Flex will treat "int" as an identifier, which is incorrect. To ensure reserved words are recognized properly, you must order your rules from **specific** (keywords) to **general** (identifiers).

## 5 Building and Testing

### 5.1 Building the Project

We use a `Makefile` to automate compilation. To build the lexer, open your terminal in the project directory and run:

```
1 make
```

This will generate the C code using Flex and compile it into an executable located in the `build/` directory.

### 5.2 Platform Specifics (Windows vs. macOS/Linux)

The provided build scripts assume a Windows environment by default (looking for `.exe` files). If you are working on **macOS** or **Linux (WSL)**, you must modify two files before starting:

1. **Makefile**: Find the line defining the executable name.
  - Change: `LEXER_EXE = $(BUILD_DIR)/lexer.exe`
  - To: `LEXER_EXE = $(BUILD_DIR)/lexer`
2. **run\_tests.py**: Find the configuration variable at the top.
  - Change: `LEXER_EXECUTABLE = "./build/lexer.exe"`
  - To: `LEXER_EXECUTABLE = "./build/lexer"`

### 5.3 Running the Test Suite

We have provided a Python script to run automated tests.

```
1 # Build first
2 make
3
4 # Run tests
5 python run_tests.py
```

The script compares your lexer's output against "Golden Output" files.

- **70%** of the grade is based on the visible tests provided to you.
- **30%** of the grade is based on **Hidden Test Cases** that check for edge cases and robustness.

## 5.4 Testing with Custom Inputs

You are encouraged to create your own test cases to debug specific issues.

1. Create a text file (e.g., `my_input.txt`) with some code.
2. Run your built lexer and redirect the input:

**On Windows:**

```
1 build\lexer.exe my_input.txt
```

**On macOS/Linux:**

```
1 ./build/lexer my_input.txt
```

This will print the tokens to the terminal, allowing you to manually verify if your rules are working as expected.

## 6 Submission Requirements

Please follow the steps below carefully to ensure your assignment is graded correctly.

### 6.1 Preparing Your Submission

Before zipping your project, you must clean the directory to remove all generated build files and executables.

1. Run the following command:

```
1 make clean  
2
```

2. Verify that the `build/` directory has been removed.

### 6.2 Naming Convention

You must compress your project folder into a single `.zip` file. The filename must strictly follow this format:

`<rollnumber>.PA1.zip`

Replace `<rollnumber>` with your actual university roll number.

**Example:**

- If your roll number is **27100289**, your submission file must be named:

`27100289_PA1.zip`